

```
# TASK 1

import cv2
import matplotlib.pyplot as plt

image_path = '/content/cvlab1img.jpg'
image = cv2.imread(image_path)

plt.figure(figsize=(10, 6))
original_image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
plt.subplot(221)
plt.imshow(original_image_rgb)
plt.title("Original Image")
plt.axis('off')

# Resize dimensions
scale_factor = 2 # Scale up by 2x
new_width = int(image.shape[1] * scale_factor)
new_height = int(image.shape[0] * scale_factor)
new_dimensions = (new_width, new_height)

# 1.1.1 Linear Interpolation
resized_linear = cv2.resize(image, new_dimensions, interpolation=cv2.INTER_LINEAR)
resized_linear_rgb = cv2.cvtColor(resized_linear, cv2.COLOR_BGR2RGB)

plt.subplot(222)
plt.imshow(resized_linear_rgb)
plt.title("Linear Interpolation")
plt.axis('off')

# 1.1.2 Nearest Neighbors Interpolation
resized_nearest = cv2.resize(image, new_dimensions, interpolation=cv2.INTER_NEAREST)
resized_nearest_rgb = cv2.cvtColor(resized_nearest, cv2.COLOR_BGR2RGB)

plt.subplot(223)
plt.imshow(resized_nearest_rgb)
plt.title("Nearest Neighbors Interpolation")
plt.axis('off')

# 1.1.3 Polynomial Interpolation (Cubic)
resized_cubic = cv2.resize(image, new_dimensions, interpolation=cv2.INTER_CUBIC)
resized_cubic_rgb = cv2.cvtColor(resized_cubic, cv2.COLOR_BGR2RGB)

plt.subplot(224)
plt.imshow(resized_cubic_rgb)
plt.title("Cubic (Polynomial) Interpolation")
plt.axis('off')

plt.tight_layout()
plt.show()
```



```
# TASK 1.2

image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

# Display original image
plt.figure(figsize=(12, 8))
plt.subplot(231)
plt.imshow(image_rgb)
plt.title("Original Image")
plt.axis('off')

# 1.2.1 Box Blurring
# Kernel size (e.g., 5x5)
box_blurred = cv2.blur(image, (5, 5))
box_blurred_rgb = cv2.cvtColor(box_blurred, cv2.COLOR_BGR2RGB)

plt.subplot(232)
plt.imshow(box_blurred_rgb)
plt.title("Box Blurring")
plt.axis('off')

# 1.2.2 Gaussian Blurring
# Kernel size (e.g., 5x5) and standard deviation
gaussian_blurred = cv2.GaussianBlur(image, (5, 5), 0)
gaussian_blurred_rgb = cv2.cvtColor(gaussian_blurred, cv2.COLOR_BGR2RGB)
```

```
plt.subplot(233)
plt.imshow(gaussian_blurred_rgb)
plt.title("Gaussian Blurring")
plt.axis('off')

# 1.2.3 Adaptive Blurring (Bilateral Filter)
# Diameter, sigmaColor, sigmaSpace (to preserve edges while blurring)
adaptive_blurred = cv2.bilateralFilter(image, 9, 75, 75)
adaptive_blurred_rgb = cv2.cvtColor(adaptive_blurred, cv2.COLOR_BGR2RGB)

plt.subplot(234)
plt.imshow(adaptive_blurred_rgb)
plt.title("Adaptive Blurring (Bilateral)")
plt.axis('off')

# Additional: Apply a stronger blur for comparison
strong_gaussian_blurred = cv2.GaussianBlur(image, (15, 15), 0)
strong_gaussian_rgb = cv2.cvtColor(strong_gaussian_blurred, cv2.COLOR_BGR2RGB)

plt.subplot(235)
plt.imshow(strong_gaussian_rgb)
plt.title("Strong Gaussian Blur")
plt.axis('off')

plt.tight_layout()
plt.show()
```



Original Image



Box Blurring



Gaussian Blurring



Adaptive Blurring (Bilateral)



Strong Gaussian Blur



```
# TASK 2
# USING DECISION TREE/RANDOM FOREST

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    confusion_matrix, roc_curve, auc, classification_report
)
from tensorflow.keras.datasets import mnist
from sklearn.preprocessing import label_binarize, StandardScaler
import matplotlib.pyplot as plt

(X_train_full, y_train_full), (X_test, y_test) = mnist.load_data()

X_train_full = X_train_full.reshape(-1, 784) / 255.0
X_test = X_test.reshape(-1, 784) / 255.0

X_train, X_val, y_train, y_val = train_test_split(X_train_full, y_train_full, test_size=0.2, random_state=42)

y_binarized = label_binarize(y_val, classes=np.unique(y_val))

model = RandomForestClassifier(n_estimators=100, random_state=42)

kf = KFold(n_splits=5, shuffle=True, random_state=42)
cv_scores = cross_val_score(model, X_train, y_train, cv=kf, scoring="accuracy")

model.fit(X_train, y_train)

y_pred = model.predict(X_test)
y_proba = model.predict_proba(X_test)

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average="macro")
```

```
recall = recall_score(y_test, y_pred, average="macro")
f1 = f1_score(y_test, y_pred, average="macro")
conf_matrix = confusion_matrix(y_test, y_pred)

fpr = {}
tpr = {}
roc_auc = {}
for i in range(10):
    fpr[i], tpr[i], _ = roc_curve(label_binarize(y_test, classes=np.unique(y_test))[:, i], y_proba[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

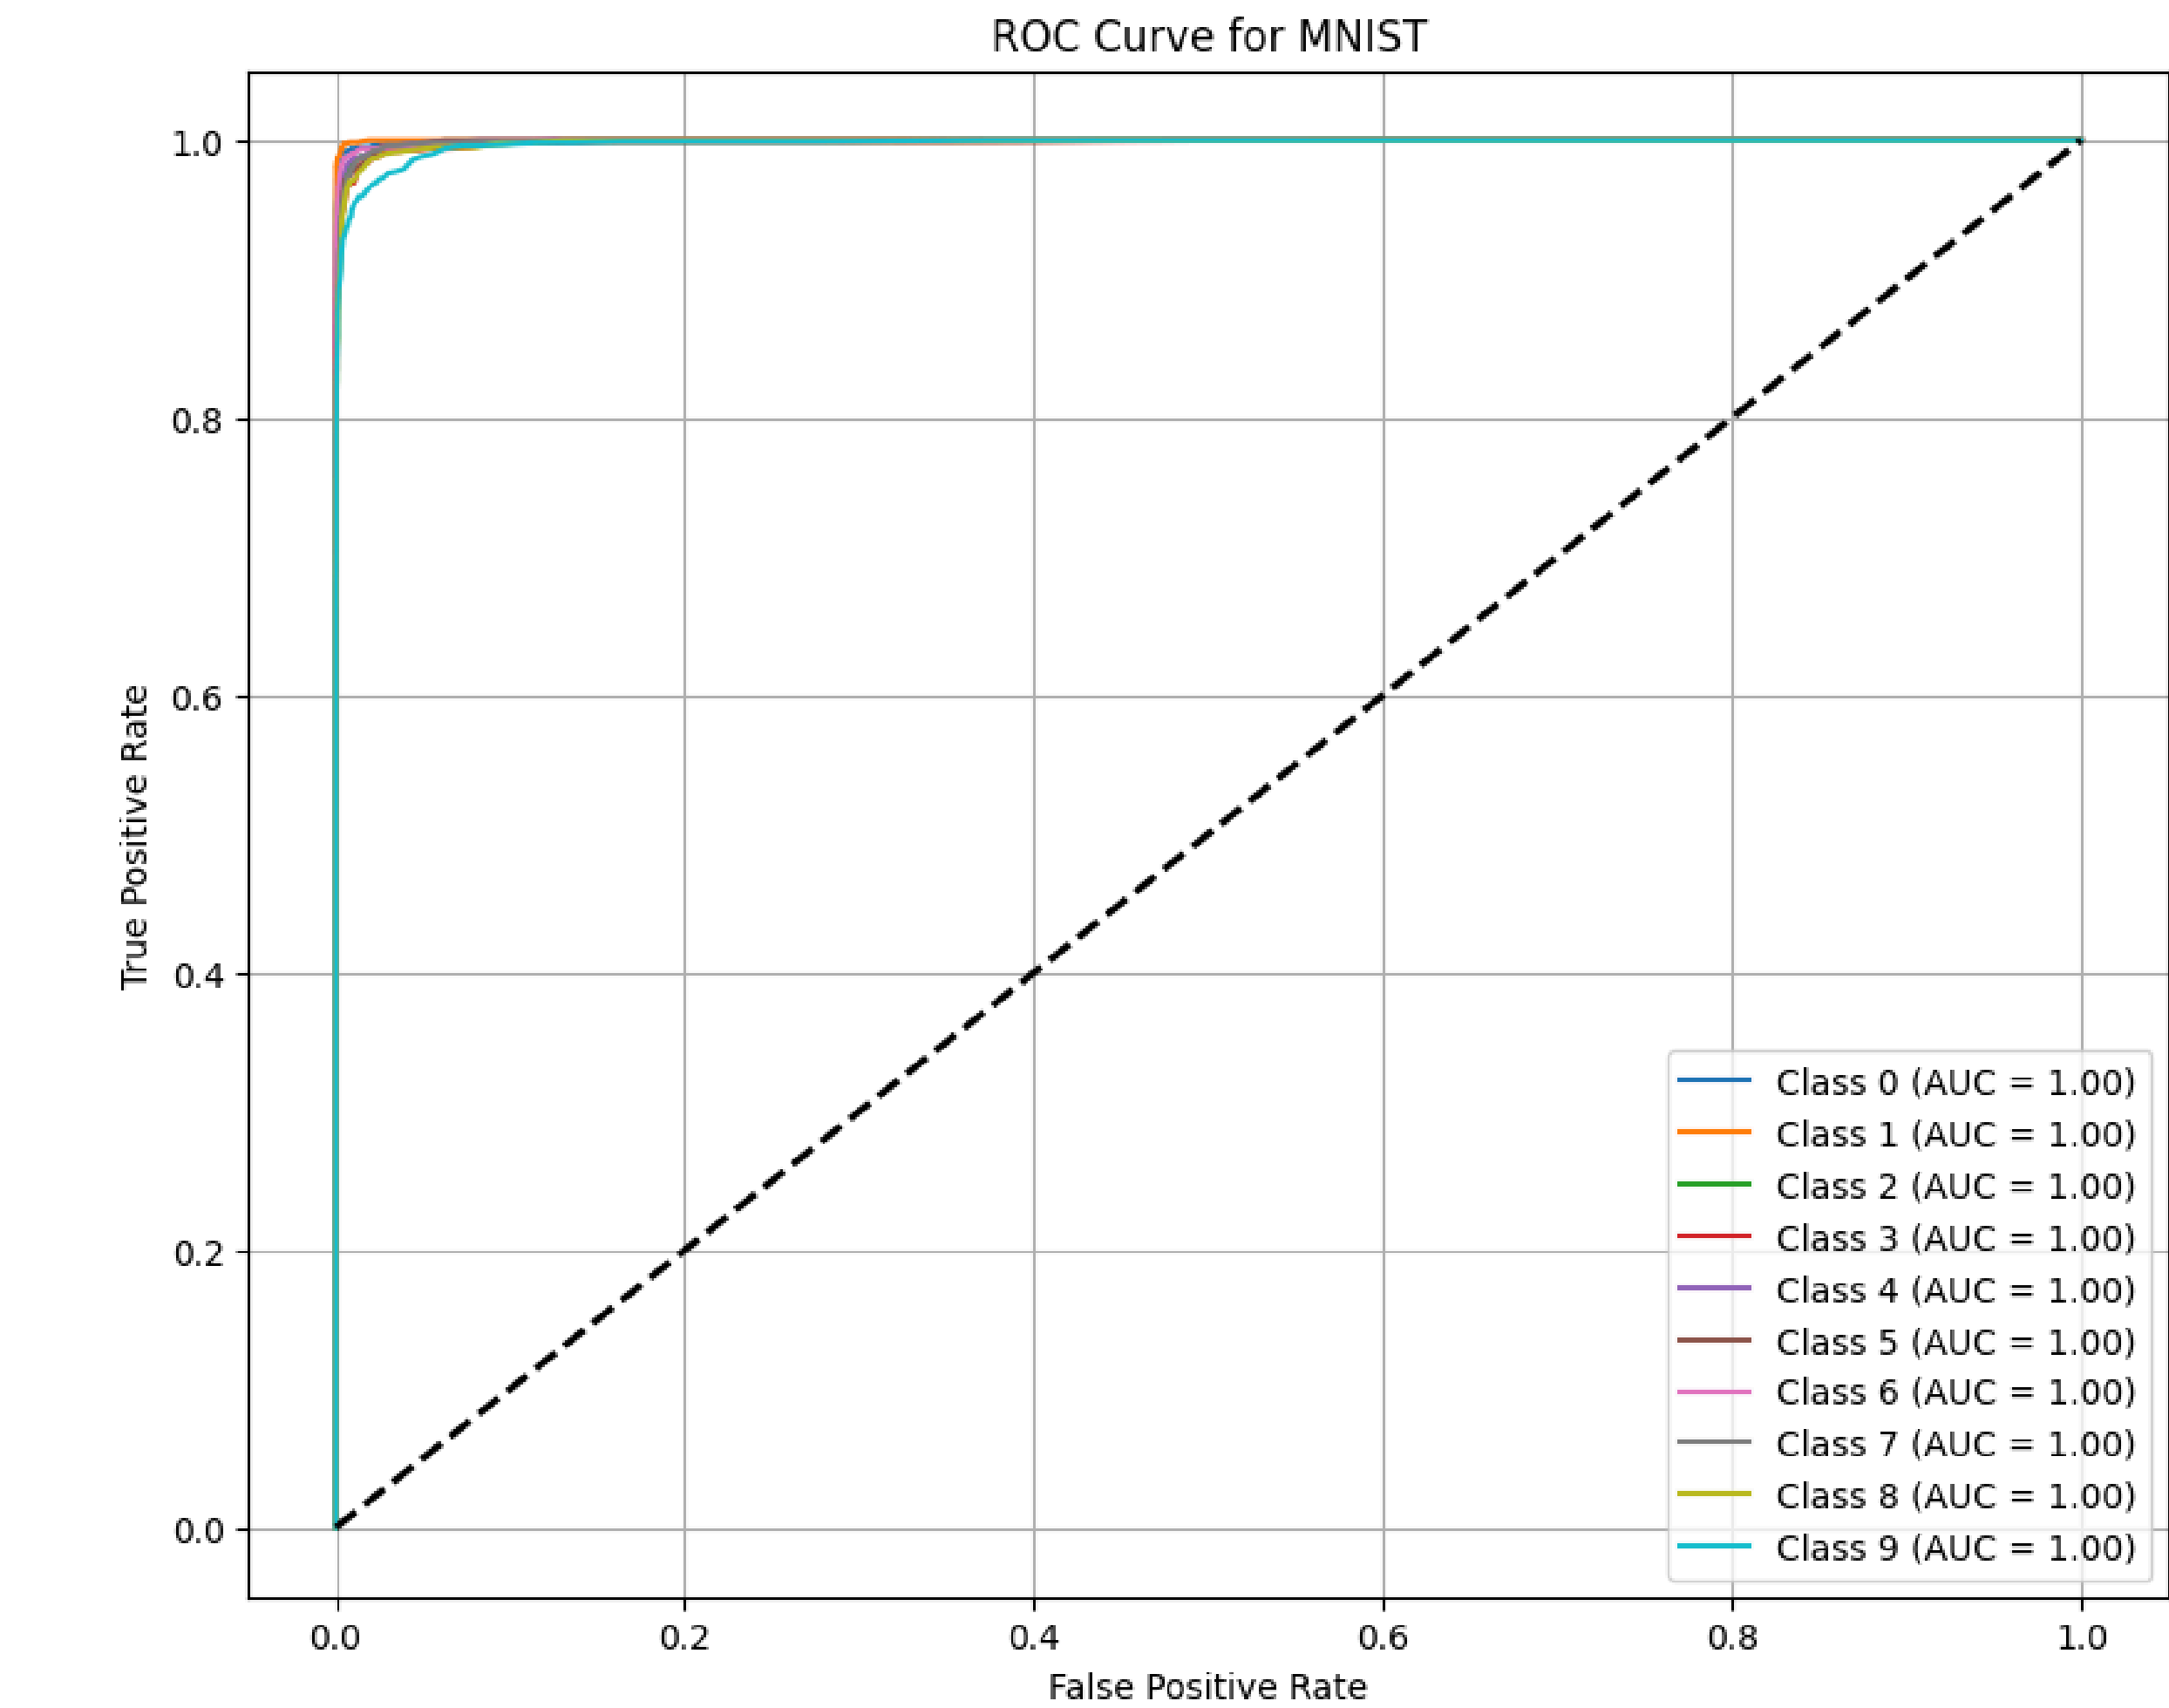
# Plot ROC Curve
plt.figure(figsize=(10, 8))
for i in range(10):
    plt.plot(fpr[i], tpr[i], label=f"Class {i} (AUC = {roc_auc[i]:.2f})")
plt.plot([0, 1], [0, 1], "k--", lw=2)
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve for MNIST")
plt.legend(loc="best")
plt.grid()
plt.show()

# Results Summary
results_summary = pd.DataFrame({
    "Metric": ["Accuracy", "Precision", "Recall", "F1-Score", "Mean CV Accuracy"],
    "Value": [accuracy, precision, recall, f1, cv_scores.mean()]
})
print("\nResults Summary:")
print(results_summary)

print("\nConfusion Matrix:")
print(conf_matrix)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

📄 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 0s 0us/step



Results Summary:										
	Metric				Value					
0	Accuracy				0.968700					
1	Precision				0.968596					
2	Recall				0.968458					
3	F1-Score				0.968500					
4	Mean CV Accuracy				0.964875					
Confusion Matrix:										
[[971	0	1	0	0	2	2	1	3	0]
[0	1123	3	2	0	2	2	1	2	0]
[6	0	998	3	4	0	4	8	9	0]
[0	0	13	971	0	5	0	9	10	2]
[1	0	3	0	956	0	4	0	2	16]
[4	0	1	11	3	857	7	1	4	4]
[8	3	0	0	2	4	938	0	3	0]
[1	3	19	4	1	0	0	985	4	11]
[3	0	2	13	5	4	2	4	931	10]
[5	6	4	10	13	3	2	5	4	957]]
Classification Report:										
	precision		recall		f1-score		support			
	0	0.97		0.99		0.98			980	
	1	0.99		0.99		0.99			1135	
	2	0.96		0.97		0.96			1032	
	3	0.96		0.96		0.96			1010	
	4	0.97		0.97		0.97			982	
	5	0.98		0.96		0.97			892	
	6	0.98		0.98		0.98			958	
	7	0.97		0.96		0.96			1028	
	8	0.96		0.96		0.96			974	
	9	0.96		0.95		0.95			1009	
	accuracy				0.97		10000			
	macro avg		0.97	0.97	0.97		10000			
	weighted avg		0.97	0.97	0.97		10000			


```
# USING ANN

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.preprocessing import StandardScaler, label_binarize

# Normalize features
scaler = StandardScaler()
X = scaler.fit_transform(X)

y_binarized = to_categorical(y)

# Split the data into train and test sets (80-20 split)
X_train, X_test, y_train, y_test = train_test_split(X, y_binarized, test_size=0.2, random_state=42)

# ANN model
def build_model():
    model = Sequential()
    model.add(Dense(64, activation='relu', input_shape=(X_train.shape[1],)))
    model.add(Dense(32, activation='relu'))
    model.add(Dense(y_binarized.shape[1], activation='softmax')) # Output layer for multi-class classification
    model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
    return model

kf = KFold(n_splits=5, shuffle=True, random_state=42)
fold accuracies = []

for train_index, val_index in kf.split(X_train):
    X_fold_train, X_fold_val = X_train[train_index], X_train[val_index]
    y_fold_train, y_fold_val = y_train[train_index], y_train[val_index]

    model = build_model()
    model.fit(X_fold_train, y_fold_train, epochs=20, batch_size=32, verbose=0)

    loss, accuracy = model.evaluate(X_fold_val, y_fold_val, verbose=0)
    fold accuracies.append(accuracy)

model = build_model()
model.fit(X_train, y_train, epochs=20, batch_size=32, verbose=0)

y_pred_proba = model.predict(X_test)
y_pred = np.argmax(y_pred_proba, axis=1)
y_test_labels = np.argmax(y_test, axis=1)

accuracy = accuracy_score(y_test_labels, y_pred)
precision = precision_score(y_test_labels, y_pred, average='macro')
recall = recall_score(y_test_labels, y_pred, average='macro')
f1 = f1_score(y_test_labels, y_pred, average='macro')
conf_matrix = confusion_matrix(y_test_labels, y_pred)

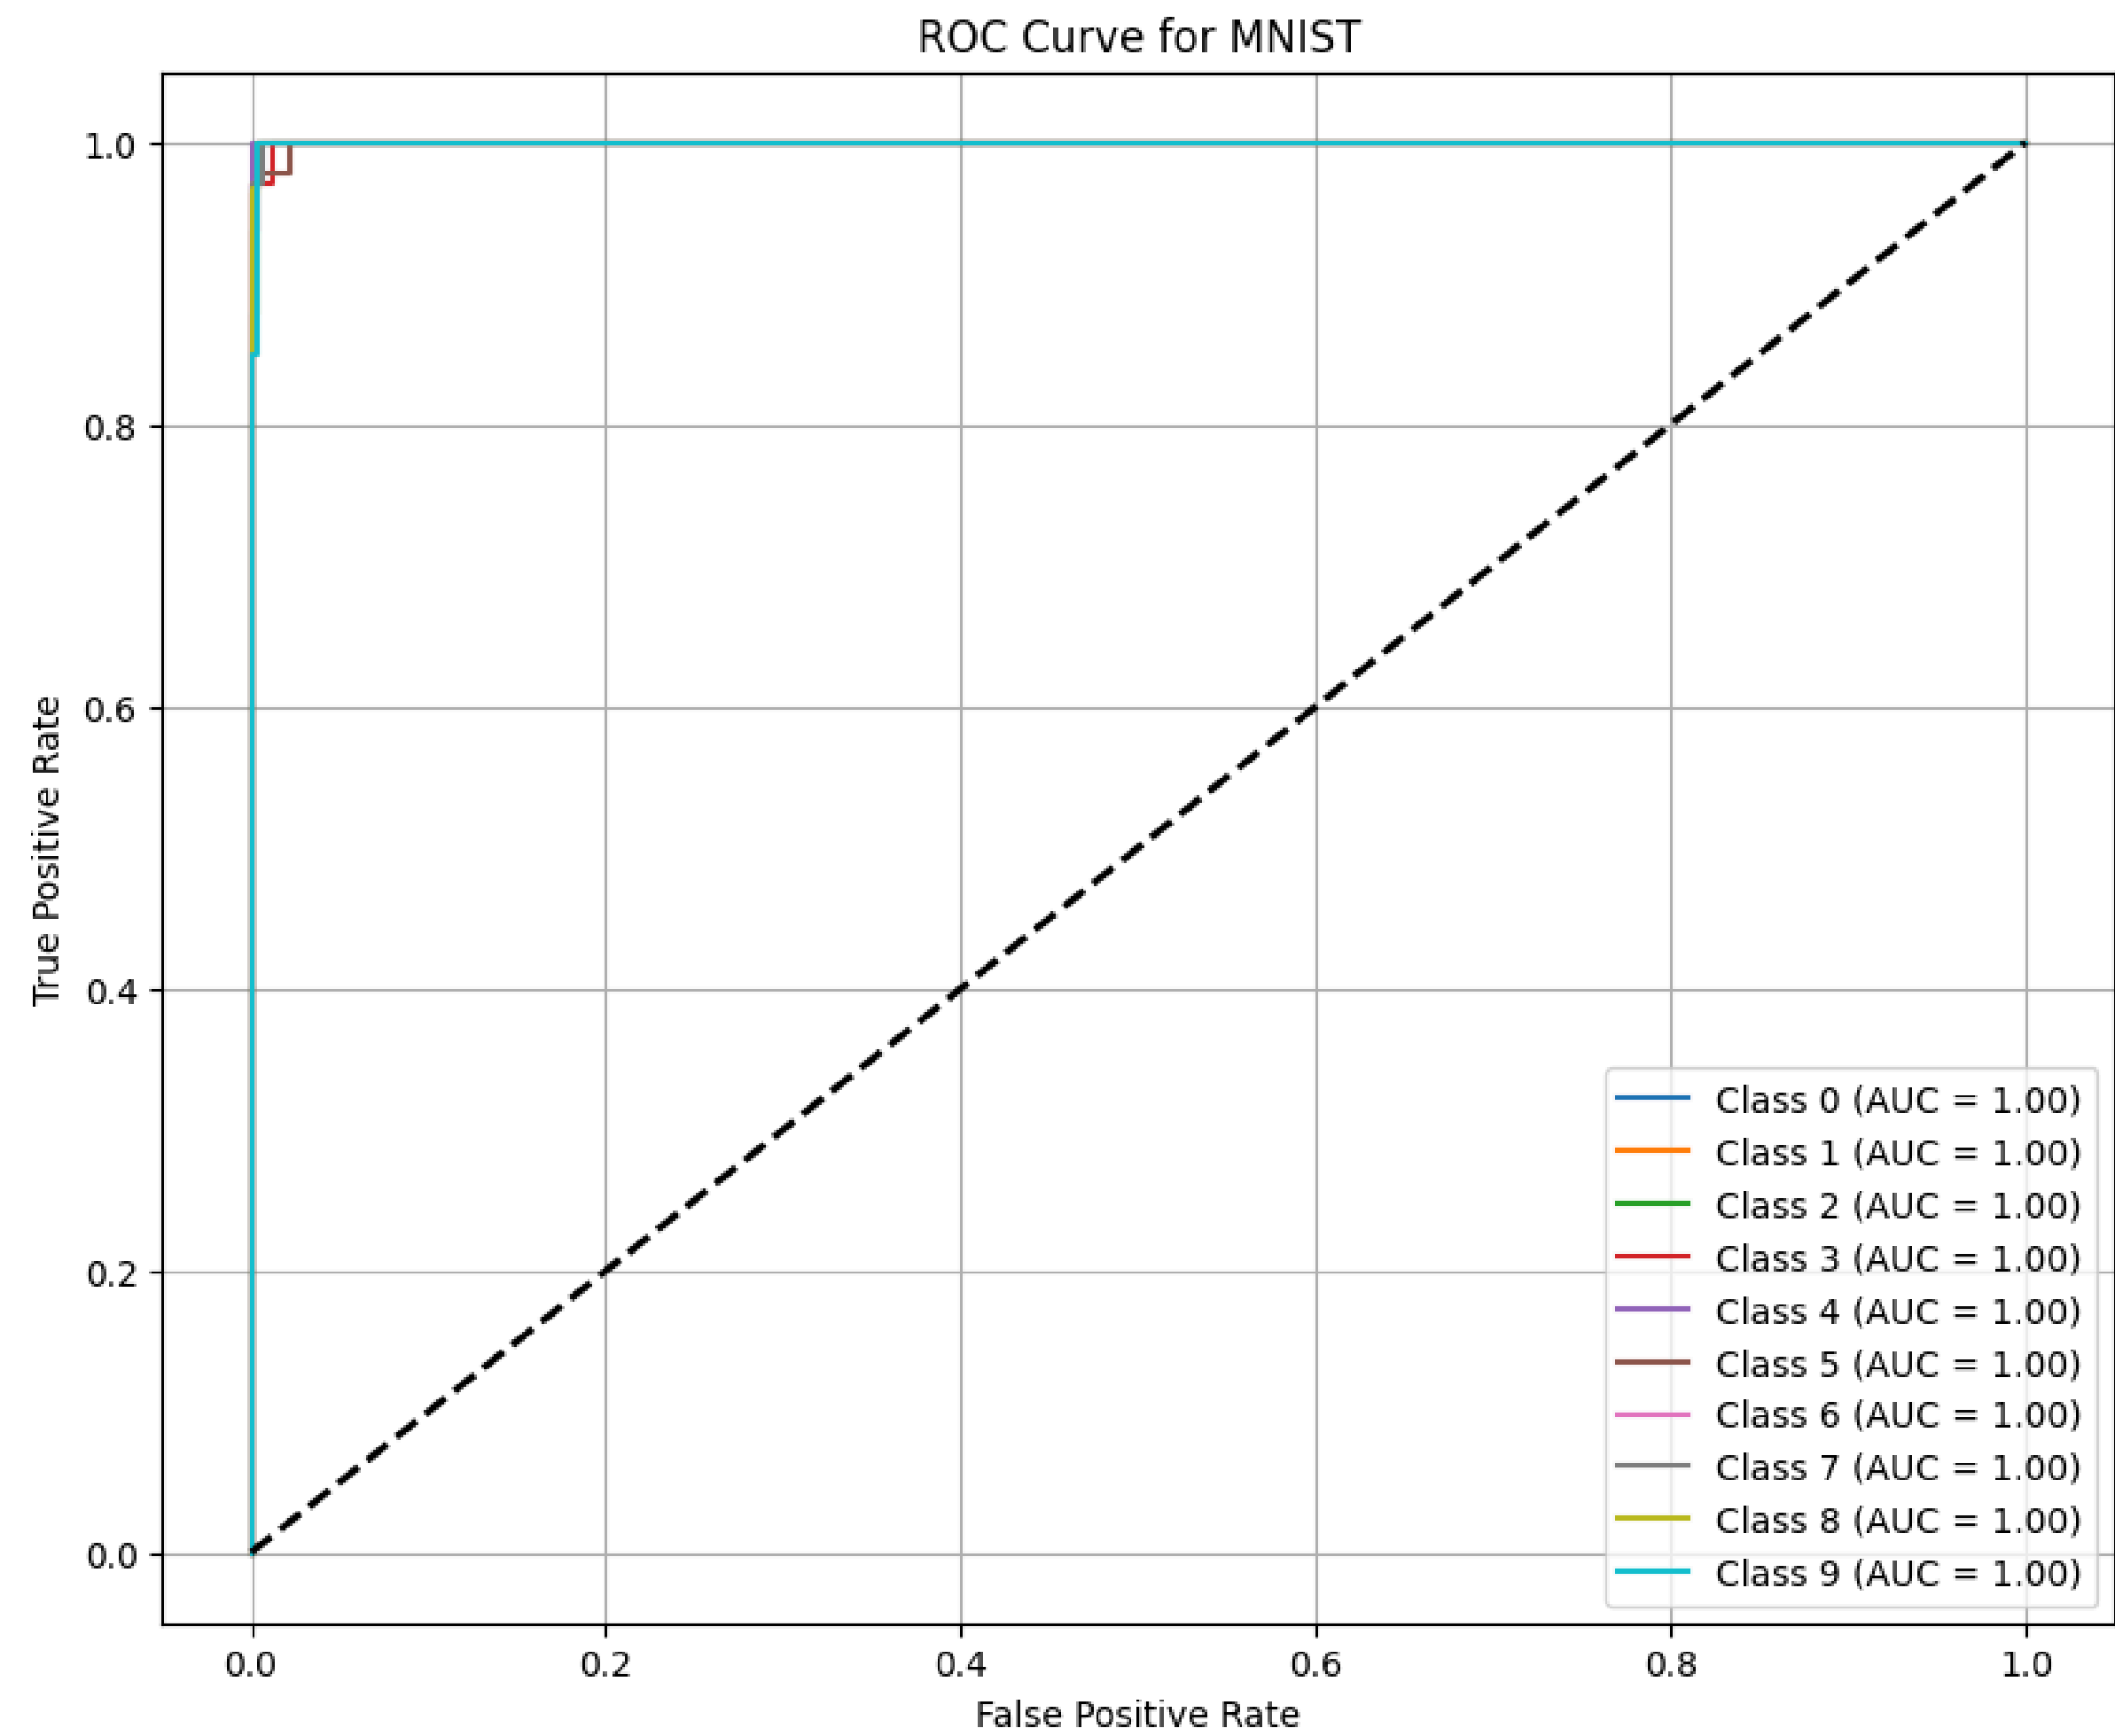
fpr = {}
tpr = {}
roc_auc = {}
for i in range(y_binarized.shape[1]):
    fpr[i], tpr[i], _ = roc_curve(y_test[:, i], y_pred_proba[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Plot ROC Curve
plt.figure(figsize=(10, 8))
for i in range(y_binarized.shape[1]):
    plt.plot(fpr[i], tpr[i], label=f"Class {i} (AUC = {roc_auc[i]:.2f})")
plt.plot([0, 1], [0, 1], "k--", lw=2)
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve for MNIST")
plt.legend(loc="best")
plt.grid()
plt.show()

# Results Summary
print("\nResults Summary:")
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")
print(f"Mean CV Accuracy: {np.mean(fold accuracies):.4f}")

print("\nConfusion Matrix:")
print(conf_matrix)

print("\nClassification Report:")
print(classification_report(y_test_labels, y_pred))
```



```
Results Summary:
Accuracy: 0.9778
Precision: 0.9779
Recall: 0.9771
F1-Score: 0.9773
Mean CV Accuracy: 0.9652
```

```
Confusion Matrix:
[[32  0  1  0  0  0  0  0  0  0]
 [ 0 27  1  0  0  0  0  0  0  0]
 [ 0  0 33  0  0  0  0  0  0  0]
 [ 0  0  0 33  0  1  0  0  0  0]
 [ 0  0  0  0 46  0  0  0  0  0]
 [ 0  0  0  0  0 45  1  0  0  1]
 [ 1  0  0  0  0  0 34  0  0  0]
 [ 0  0  0  0  0  0  0 33  0  1]
 [ 0  1  0  0  0  0  0  0 29  0]
 [ 0  0  0  0  0  0  0  0  0 40]]
```

Classification Report:				
	precision	recall	f1-score	support
0	0.97	0.97	0.97	33
1	0.96	0.96	0.96	28
2	0.94	1.00	0.97	33
3	1.00	0.97	0.99	34
4	1.00	1.00	1.00	46
5	0.98	0.96	0.97	47
6	0.97	0.97	0.97	35
7	1.00	0.97	0.99	34
8	1.00	0.97	0.98	30
9	0.95	1.00	0.98	40
accuracy			0.98	360
macro avg	0.98	0.98	0.98	360
weighted avg	0.98	0.98	0.98	360