

```
import cv2
import numpy as np
from PIL import Image
import io
import os

from google.colab import files
uploaded = files.upload()

filename = list(uploaded.keys())[0]

image = cv2.imread(filename)

image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

original_size = os.path.getsize(filename) / 1024
print(f"Original Image Size: {original_size:.2f} KB")

jpeg_filename = "compressed_image.jpg"
cv2.imwrite(jpeg_filename, image, [cv2.IMWRITE_JPEG_QUALITY, 30])

jpeg_size = os.path.getsize(jpeg_filename) / 1024
print(f"JPEG Compressed Size: {jpeg_size:.2f} KB")

png_filename = "compressed_image.png"
cv2.imwrite(png_filename, image, [cv2.IMWRITE_PNG_COMPRESSION, 9])

png_size = os.path.getsize(png_filename) / 1024
print(f"PNG Compressed Size: {png_size:.2f} KB")

import matplotlib.pyplot as plt

fig, ax = plt.subplots(1, 3, figsize=(15,5))

ax[0].imshow(image_rgb)
ax[0].set_title("Original Image")
ax[0].axis("off")

jpeg_img = Image.open(jpeg_filename)
ax[1].imshow(jpeg_img)
ax[1].set_title("JPEG Compressed Image")
ax[1].axis("off")

png_img = Image.open(png_filename)
ax[2].imshow(png_img)
ax[2].set_title("PNG Compressed Image")
ax[2].axis("off")

plt.show()
```



Choose Files nature-4k-p...-preview.jpg

- **nature-4k-pc-full-hd-wallpaper-preview.jpg**(image/jpeg) - 104758 bytes, last modified: 2/14/2025 - 100% done

Saving nature-4k-pc-full-hd-wallpaper-preview.jpg to nature-4k-pc-full-hd-wallpaper-preview (1).jpg
 Original Image Size: 102.30 KB
 JPEG Compressed Size: 40.21 KB
 PNG Compressed Size: 587.04 KB

Original Image



JPEG Compressed Image



PNG Compressed Image



```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from sklearn.metrics import classification_report, roc_auc_score, roc_curve, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
# Load MNIST dataset
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
```

```
# Normalize pixel values
x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
# Reshape to match CNN input (28x28 images with 1 channel)
x_train = x_train.reshape(-1, 28, 28, 1)
x_test = x_test.reshape(-1, 28, 28, 1)
```

```
# One-hot encode labels
y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)
```



Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
 11490434/11490434 ————— 0s 0us/step

```
# Define CNN model
model = keras.Sequential([
    layers.Conv2D(32, (3,3), strides=2, activation='relu', input_shape=(28,28,1)),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64, (3,3), strides=2, activation='relu'),
    layers.MaxPooling2D((2,2)),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(10, activation='softmax') # 10 classes
])
```

```
# Compile the model
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

```
history = model.fit(x_train, y_train, epochs=3, batch_size=128, validation_split=0.2)
```



```
/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `i
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Epoch 1/3
375/375 ————— 9s 19ms/step - accuracy: 0.6667 - loss: 1.1541 - val_accuracy: 0.9377 - val_loss: 0.2090
Epoch 2/3
375/375 ————— 6s 16ms/step - accuracy: 0.9413 - loss: 0.1948 - val_accuracy: 0.9557 - val_loss: 0.1491
Epoch 3/3
375/375 ————— 10s 15ms/step - accuracy: 0.9570 - loss: 0.1456 - val_accuracy: 0.9607 - val_loss: 0.1342
```


```
# Predictions
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = np.argmax(y_test, axis=1)

# Model Evaluation
accuracy = np.mean(y_pred_classes == y_true)
precision = classification_report(y_true, y_pred_classes, output_dict=True)['weighted avg']['precision']
recall = classification_report(y_true, y_pred_classes, output_dict=True)['weighted avg']['recall']
f1_score = classification_report(y_true, y_pred_classes, output_dict=True)['weighted avg']['f1-score']

# Confusion Matrix
conf_matrix = confusion_matrix(y_true, y_pred_classes)
plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap="Blues")
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix for MNIST')
plt.show()

# ROC Curve & AUC
auc_score = roc_auc_score(y_test, y_pred, multi_class='ovr')
print(f"AUC Score: {auc_score:.4f}")
```

313/313 — 1s 2ms/step



AUC Score: 0.9999

```
# Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()

# Normalize pixel values
x_train, x_test = x_train / 255.0, x_test / 255.0

# One-hot encode labels
y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)

# Define CNN model
model = keras.Sequential([
```

```

layers.Conv2D(32, (3,3), activation='relu', input_shape=(32,32,3)),
layers.MaxPooling2D((2,2)),
layers.Conv2D(64, (3,3), activation='relu'),
layers.MaxPooling2D((2,2)),
layers.Conv2D(128, (3,3), activation='relu'),
layers.MaxPooling2D((2,2)),
layers.Flatten(),
layers.Dense(128, activation='relu'),
layers.Dense(10, activation='softmax') # 10 classes
])

# Compile the model
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(x_train, y_train, epochs=20, batch_size=128, validation_split=0.2)

```



Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>

170498071/170498071 ————— 4s 0us/step

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `i`
super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Epoch 1/20

313/313 ————— 52s 160ms/step - accuracy: 0.2894 - loss: 1.9187 - val_accuracy: 0.4972 - val_loss: 1.4064

Epoch 2/20

313/313 ————— 82s 161ms/step - accuracy: 0.5148 - loss: 1.3564 - val_accuracy: 0.5576 - val_loss: 1.2370

Epoch 3/20

313/313 ————— 79s 154ms/step - accuracy: 0.5816 - loss: 1.1830 - val_accuracy: 0.5931 - val_loss: 1.1507

Epoch 4/20

313/313 ————— 84s 160ms/step - accuracy: 0.6282 - loss: 1.0615 - val_accuracy: 0.6238 - val_loss: 1.0842

Epoch 5/20

313/313 ————— 48s 152ms/step - accuracy: 0.6560 - loss: 0.9841 - val_accuracy: 0.6532 - val_loss: 1.0039

Epoch 6/20

313/313 ————— 48s 152ms/step - accuracy: 0.6858 - loss: 0.9070 - val_accuracy: 0.6708 - val_loss: 0.9650

Epoch 7/20

313/313 ————— 85s 163ms/step - accuracy: 0.7027 - loss: 0.8554 - val_accuracy: 0.6819 - val_loss: 0.9211

Epoch 8/20

313/313 ————— 48s 154ms/step - accuracy: 0.7200 - loss: 0.8094 - val_accuracy: 0.7000 - val_loss: 0.8831

Epoch 9/20

313/313 ————— 81s 152ms/step - accuracy: 0.7391 - loss: 0.7529 - val_accuracy: 0.6942 - val_loss: 0.9102

Epoch 10/20

313/313 ————— 85s 161ms/step - accuracy: 0.7520 - loss: 0.7138 - val_accuracy: 0.6905 - val_loss: 0.9209

Epoch 11/20

313/313 ————— 48s 154ms/step - accuracy: 0.7661 - loss: 0.6711 - val_accuracy: 0.7055 - val_loss: 0.8788

Epoch 12/20

313/313 ————— 84s 161ms/step - accuracy: 0.7836 - loss: 0.6260 - val_accuracy: 0.7003 - val_loss: 0.9017

Epoch 13/20

313/313 ————— 79s 151ms/step - accuracy: 0.7909 - loss: 0.5969 - val_accuracy: 0.7126 - val_loss: 0.8844

Epoch 14/20

313/313 ————— 51s 161ms/step - accuracy: 0.8083 - loss: 0.5510 - val_accuracy: 0.7075 - val_loss: 0.8844

Epoch 15/20

313/313 ————— 82s 161ms/step - accuracy: 0.8144 - loss: 0.5265 - val_accuracy: 0.7084 - val_loss: 0.9089

Epoch 16/20

313/313 ————— 81s 158ms/step - accuracy: 0.8278 - loss: 0.4946 - val_accuracy: 0.7159 - val_loss: 0.9150

Epoch 17/20

313/313 ————— 81s 153ms/step - accuracy: 0.8369 - loss: 0.4686 - val_accuracy: 0.7137 - val_loss: 0.9149

Epoch 18/20

313/313 ————— 51s 162ms/step - accuracy: 0.8496 - loss: 0.4281 - val_accuracy: 0.7185 - val_loss: 0.9235

Epoch 19/20

313/313 ————— 79s 154ms/step - accuracy: 0.8638 - loss: 0.3958 - val_accuracy: 0.7150 - val_loss: 0.9577

Epoch 20/20

313/313 ————— 50s 158ms/step - accuracy: 0.8718 - loss: 0.3699 - val_accuracy: 0.6989 - val_loss: 1.0282

```

y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = np.argmax(y_test, axis=1)

```

```

accuracy = np.mean(y_pred_classes == y_true)
precision = classification_report(y_true, y_pred_classes, output_dict=True)['weighted avg']['precision']
recall = classification_report(y_true, y_pred_classes, output_dict=True)['weighted avg']['recall']
f1_score = classification_report(y_true, y_pred_classes, output_dict=True)['weighted avg']['f1-score']

```

```
conf_matrix = confusion_matrix(y_true, y_pred_classes)
```

```

conf_matrix = confusion_matrix(y_true, y_pred_classes)
plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap="Blues")
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix for CIFAR-10')
plt.show()


```

```

auc_score = roc_auc_score(y_test, y_pred, multi_class='ovr')
print(f"AUC Score: {auc_score:.4f}")

```

313/313 ————— 5s 14ms/step



AUC Score: 0.8520